

COMP3242/6242: Deep Learning

Week 2 — Linear Classifiers and Multilayer Perceptrons

Stephen Gould
stephen.gould@anu.edu.au

Australian National University

Semester 1, 2026

Overview


- ▶ binary classification
- ▶ logistic regression
- ▶ the XOR problem
- ▶ multilayer perceptron
- ▶ activation functions
- ▶ the universal approximation theorem
- ▶ multi-class classification
- ▶ loss functions
- ▶ gradient descent optimization
- ▶ calculating gradients

Binary (linear) Classification


- ▶ **data:** training examples $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^N$ composed of
 - ▶ features $x^{(i)} \in \mathbb{R}^n$
 - ▶ target labels $y^{(i)} \in \{0, 1\}$
- ▶ **regression task:** learn a density function $P(y | x)$ from the training examples, which estimates the probability of 0 or 1 on an input x
- ▶ **classifier task:** learn a linear classifier $f(x) = a^T x + b$ from the training examples, which predicts 1 on an input x if $a^T x + b \geq 0$ and 0 otherwise

2D Illustration

separable case:



inseparable case:



- ▶ $a^T x + b$ measures (signed) distance from the hyperplane

Logistic Regression

Let $y \in \{0, 1\}$ be a random variable with distribution

$$P(y = 1 \mid x) = \frac{\exp(a^T x + b)}{1 + \exp(a^T x + b)}$$

where a, b are parameters and $x \in \mathbb{R}^n$ is observed features.

The **maximum likelihood principle** states that we should choose parameters a, b that maximizes the probability of the observed data under the model,


$$P(\mathcal{D}) = \prod_{i=1}^N P(y^{(i)} \mid x^{(i)})$$

In practice we (equivalently) minimize the negative log-likelihood function instead,

$$\begin{aligned}\ell(a, b; \mathcal{D}) &= -\log \left(\prod_{i:y^{(i)}=1} \frac{\exp(a^T x^{(i)} + b)}{1 + \exp(a^T x^{(i)} + b)} \prod_{i:y^{(i)}=0} \frac{1}{1 + \exp(a^T x^{(i)} + b)} \right) \\ &= - \sum_{i=1}^N (a^T x^{(i)} + b) + \sum_{i=1}^N \log(1 + \exp(a^T x^{(i)} + b))\end{aligned}$$

Logistic Function


$$y = \sigma(z) = \frac{\exp(z)}{\exp(z) + 1} = \frac{1}{1 + \exp(-z)}$$







- ▶ observe that $\sigma : \mathbb{R} \rightarrow [0, 1]$ and $\sigma(z) \geq \frac{1}{2} \iff z = a^T x + b \geq 0$

Logistic Function (2D Feature Space)

$$y = \frac{1}{1 + \exp(-a^T x - b)}$$



2D Illustration Revisited



Regularization



- ▶ we often don't want our classifiers to be too confident in the margin between positive and negative samples
- ▶ regularization is one way to reduce confidence in that region and generalize better to unseen data
- ▶ add a penalty on the magnitude of parameters

$$\text{minimize}_{a,b} \quad \ell(a, b; \mathcal{D}) + \frac{\lambda}{2} \|a\|^2$$

- ▶ also called weight decay

$$\nabla_a \left(\ell(a, b; \mathcal{D}) + \frac{\lambda}{2} \|a\|^2 \right) = \nabla_a \ell + \lambda a$$


- ▶ we'll see other techniques for preventing over fitting later, e.g., data augmentation



Temperature Scaling

Introduce temperature parameter $\tau > 0$,


$$y = \sigma\left(\frac{z}{\tau}\right) = \frac{1}{1 + \exp(-z/\tau)}$$




- ▶ high τ has similar affect to regularization (restricted to logistic and not trained)
- ▶ can be applied after training to sharpen/flatten distribution

Logistic Regression as a Single Neuron


classic view, $y = \sigma(\sum_{i=1}^n a_i x_i + b)$



more modern view, $y = \sigma(a^T x + b)$




XOR Problem




- ▶ cannot be separated by a (single) linear decision boundary
- ▶ but can be separated by an additional layer of processing

XOR Problem: Two Layers


NAND



OR



AND



(the first layer defines a new features space; here we are visualizing in the original feature space)


Multilayer Perceptron

(Rosenblatt 1961, Widrow & Hoff 1960)

Compose multiple logistic layers together. E.g., for two-layers,

$$y = \sigma(c^T \sigma(Ax + b) + d)$$


where σ is applied elementwise. Here $A \in \mathbb{R}^{m \times n}$, $b \in \mathbb{R}^m$, $c \in \mathbb{R}^m$, and $d \in \mathbb{R}$.



Activation Functions


- ▶ logistic function (sigmoid)

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$




- ▶ hyperbolic tangent

$$\sigma(z) = \tanh(z) = \frac{e^{2z} - 1}{e^{2z} + 1}$$



- ▶ rectified linear unit (ReLU)

$$\sigma(z) = \max\{0, z\}$$



- ▶ and many others (e.g., Leaky ReLU, GLU, etc.)

Universal Approximation Theorem

(Cybenko 1989, Hornik 1991)

Theorem (roughly):


- ▶ let $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ be a (well-behaved) activation function
- ▶ let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be any continuous function
- ▶ then there exists parameters $A \in \mathbb{R}^{m \times n}$, $b \in \mathbb{R}^m$, $c \in \mathbb{R}^m$, and $d \in \mathbb{R}$ such that the function

$$\hat{f}(x) = c^T \sigma(Ax + b) + d$$

approximates $f(x)$ everywhere

- ▶ that is, $|\hat{f}(x) - f(x)| \leq \epsilon$ for all x

Proof Sketch (1D)




Why Deep Learning?

- ▶ so we can approximate any function with a two-layer network
- ▶ **problem 1:** doesn't tell us how many parameters
 - ▶ still active area of research, but indication is that we need $O(e^{1/\epsilon})$, in general
- ▶ **problem 2:** doesn't tell us how to find the parameters efficiently
- ▶ we can reduce the number of parameters by going deeper
- ▶ learning turns out to be “easier” in deeper networks
 - ▶ currently not well understood
- ▶ in practice we choose activation functions that violate the conditions of the theorem, e.g., ReLU

Advanced Topic: Identifiability

- ▶ consider the two-layer perceptron (w/ or w/o second activation function),



- ▶ let $P \in \mathbb{R}^{m \times m}$ be a permutation matrix, then

$$\begin{aligned} y &= c^T \sigma(Ax + b) + d \\ &= c^T \sigma(P^{-1}PAx + P^{-1}Pb) + d && (\text{since } P^{-1}P = I) \\ &= c^T P^{-1} \sigma(PAx + Pb) + d && (\text{since } \sigma \text{ is applied elementwise}) \\ &= \tilde{c}^T \sigma(\tilde{A}x + \tilde{b}) + d \end{aligned}$$

- ▶ so neural network parameters are never unique

Multi-class Classification

- ▶ **task:** K -way classification (e.g., ImageNet)
- ▶ **data:** training examples $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^N$ composed of
 - ▶ features $x^{(i)} \in \mathbb{R}^n$
 - ▶ target labels $y^{(i)} \in \{1, \dots, K\}$
- ▶ we can extend the logistic function to the multi-class logistic

$$P(y = k \mid x) = \frac{\exp(a_k^T x + b_k)}{\sum_{j=1}^K \exp(a_j^T x + b_j)} \quad \text{for } k = 1, \dots, K$$

parametrized by $\{(a_k, b_k) \mid k = 1, \dots, K\}$

- ▶ the function that takes vector $z = (z_1, \dots, z_K)$ and computes $(\frac{\exp z_1}{Z}, \dots, \frac{\exp z_K}{Z})$ with $Z = \sum_{k=1}^K \exp z_k$ is called **softmax**; the z_k are called **logits**
- ▶ can be computed in a numerically stable way (how?)

Multi-label Classification

- ▶ **task:** set of K binary classification problems (e.g., attribute classification)
- ▶ **data:** training examples $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^N$ composed of
 - ▶ features $x^{(i)} \in \mathbb{R}^n$
 - ▶ target labels $y^{(i)} \in \{0, 1\}^K$
- ▶ model with a set of K sigmoid functions

$$P(y_k = 1 \mid x) = \frac{\exp(a_k^T x + b_k)}{1 + \exp(a_k^T x + b_k)} \quad \text{for } k = 1, \dots, K$$



parametrized by $\{(a_k, b_k) \mid k = 1, \dots, K\}$

As Multilayer Perceptrons

with parameters $A \in \mathbb{R}^{m \times n}$, $b \in \mathbb{R}^m$, $C \in \mathbb{R}^{K \times m}$, and $d \in \mathbb{R}^K$,

$$y^{\text{multi-class}} = \text{softmax}(C\sigma(Ax + b) + d)$$

$$y^{\text{multi-label}} = \text{sigmoid}(C\sigma(Ax + b) + d)$$



Common Loss Functions

- ▶ loss functions tell the optimizer what to do
- ▶ a surrogate for what we really care about
- ▶ typically decompose over training examples, $(x, y) \sim \mathcal{D}$

Mean Square Error. standard loss for regression, $\frac{1}{2} \|f_\theta(x) - y\|^2$

0-1 Loss. what we care about for classification, 0 if $\llbracket f_\theta(x) \geq 0 \rrbracket = y$ and 1 otherwise
but 0-1 loss is non-differentiable so difficult to optimize

Negative Log-likelihood. standard loss for classification, $-\log p_\theta(y | x)$


Cross Entropy. negative log-likelihood loss where p_θ is the multi-class logistic,

$$p_\theta \propto \exp(f_\theta(x))$$

Binary Cross Entropy. negative log-likelihood loss for K independent binary variables,

$$-\sum_{k=1}^K y_k \log p_{\theta,k}(1 | x) + (1 - y_k) \log p_{\theta,k}(0 | x)$$

Gradient Descent Optimization



- ▶ function decreases in the negative gradient direction
- ▶ but if we move too far the function may increase again

Multi-variate Calculus Review: Gradients

Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be some function, fix some $x \in \mathbb{R}^n$, and consider the expression

$$\lim_{\alpha \rightarrow 0} \frac{f(x + \alpha e_i) - f(x)}{\alpha}$$

where $e_i = (0, \dots, 1, 0, \dots)$ is the i -th canonical vector.

If the limit exists, it's called the i -th **partial derivative** of f at x and denoted by $\frac{\partial f(x)}{\partial x_i}$.

The **gradient** of $f(x)$ with respect to $x \in \mathbb{R}^n$ is the vector of partial derivatives

$$\nabla_x f(x) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$$

Gradient Descent Algorithm

$$\theta^{(t+1)} = \theta^{(t)} - \eta \nabla \ell(\theta^{(t)})$$

- ▶ η is the (iteration dependent) step size or step length or learning rate
 - ▶ ideally chosen by line search (expensive)
 - ▶ instead, in deep learning, chosen by fixed schedule

```
1 def gradient_descent(loss, theta0):
2     """Generic gradient descent method."""
3
4     theta = theta0
5     for t in range(max_iters):
6         dtheta = -1.0 * gradient(loss, theta)    # descent direction is negative of gradient at theta
7         eta = line_search(loss, theta, dtheta)    # choose step size by line search
8         theta = theta + eta * dtheta              # update
9
10        if (converged()):                      # check stopping criteria
11            break
12
13    return theta
```

- ▶ stopping criterion usually of the form $\|\nabla \ell(\theta)\|_2 \leq \epsilon$
- ▶ very simple—just need gradient of ℓ with respect to θ —but can be very slow

(Vanilla) Gradient Descent Optimization in PyTorch

```
1  dataloader = ...      # way of loading batches of input-label pairs for training
2  model = ...          # definition of our prediction function
3
4  criterion = torch.nn.CrossEntropyLoss(reduction='mean')
5  optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
6
7  for epoch in range(max_epochs):
8      for iteration, batch in enumerate(dataloader, 0):
9          inputs, labels = batch
10
11         # zero gradient buffers
12         optimizer.zero_grad()
13
14         # compute model outputs for given inputs
15         outputs = model(inputs)
16
17         # compute and print the loss
18         loss = criterion(outputs, labels)
19         print(loss.item())
20
21         # compute gradient of the loss wrt model parameters
22         loss.backward()
23
24         # take a gradient step
25         optimizer.step()
```


Example Quadratic Problem in \mathbb{R}^2

(Boyd and Vandenberghe, 2024)

$$\ell(\theta) = \frac{1}{2} (\theta_1^2 + \gamma \theta_2^2) \quad (\gamma > 0)$$

with exact line search, starting at $\theta^{(0)} = (\gamma, 1)$:

$$\theta_1^{(t)} = \gamma \left(\frac{\gamma - 1}{\gamma + 1} \right)^t, \quad \theta_2^{(t)} = \left(-\frac{\gamma - 1}{\gamma + 1} \right)^t$$



- ▶ very slow if $\gamma \gg 1$ or $\gamma \ll 1$
- ▶ example above is for $\gamma = 10$

Calculating Gradients: Worked Example

- ▶ two-layer perceptron with sigmoid activation

$$z = \sigma(Ax + b)$$

$$y = \sigma(c^T z + d)$$

where $\sigma(\xi) = (1 + e^{-\xi})^{-1}$ is applied elementwise

- ▶ arbitrary loss function to minimize ℓ
- ▶ for gradient descent we need $\frac{d\ell}{dA}$, $\frac{d\ell}{db}$, etc.
- ▶ observe that

$$\begin{aligned}\frac{d\sigma}{d\xi} &= (-e^{-\xi})(-1)(1 + e^{-\xi})^{-2} \\ &= \left(\frac{1}{1 + e^{-\xi}} \right) \left(\frac{e^{-\xi}}{1 + e^{-\xi}} \right) \\ &= \sigma(\xi)(1 - \sigma(\xi))\end{aligned}$$

Calculating Gradients: Worked Example (cont.)

- ▶ let $\xi = c^T z + d$
- ▶ layer 2 parameters:

$$\begin{aligned}\frac{\partial \ell}{\partial d} &= \frac{d\ell}{dy} \frac{\partial y}{\partial d} \\ &= \frac{d\ell}{dy} \frac{dy}{d\xi} \frac{\partial \xi}{\partial d} \\ &= \frac{d\ell}{dy} y(1 - y)\end{aligned}$$

$$\begin{aligned}\nabla_c \ell &= \frac{d\ell}{dy} \nabla_c y \\ &= \frac{d\ell}{dy} \frac{dy}{d\xi} \nabla_c \xi \\ &= \frac{d\ell}{dy} y(1 - y)z\end{aligned}$$

- ▶ where the last line is since $y = \sigma(\xi)$

Calculating Gradients: Worked Example (cont.)

- ▶ layer 2 input / layer 1 output:

$$\nabla_z \ell = \frac{d\ell}{dy} y(1-y)c$$

- ▶ layer 1 parameters:

$$\nabla_b \ell = \nabla_z \ell \circ z \circ (1-z)$$

$$\nabla_A \ell = \frac{d\ell}{dy} y(1-y) (c \circ z \circ (1-z)) x^T$$

where \circ denotes elementwise product

- ▶ note that we have reused calculations from evaluation of ℓ ; we could have written y and z in terms of x in the gradient expressions (memory-vs-compute trade-off)